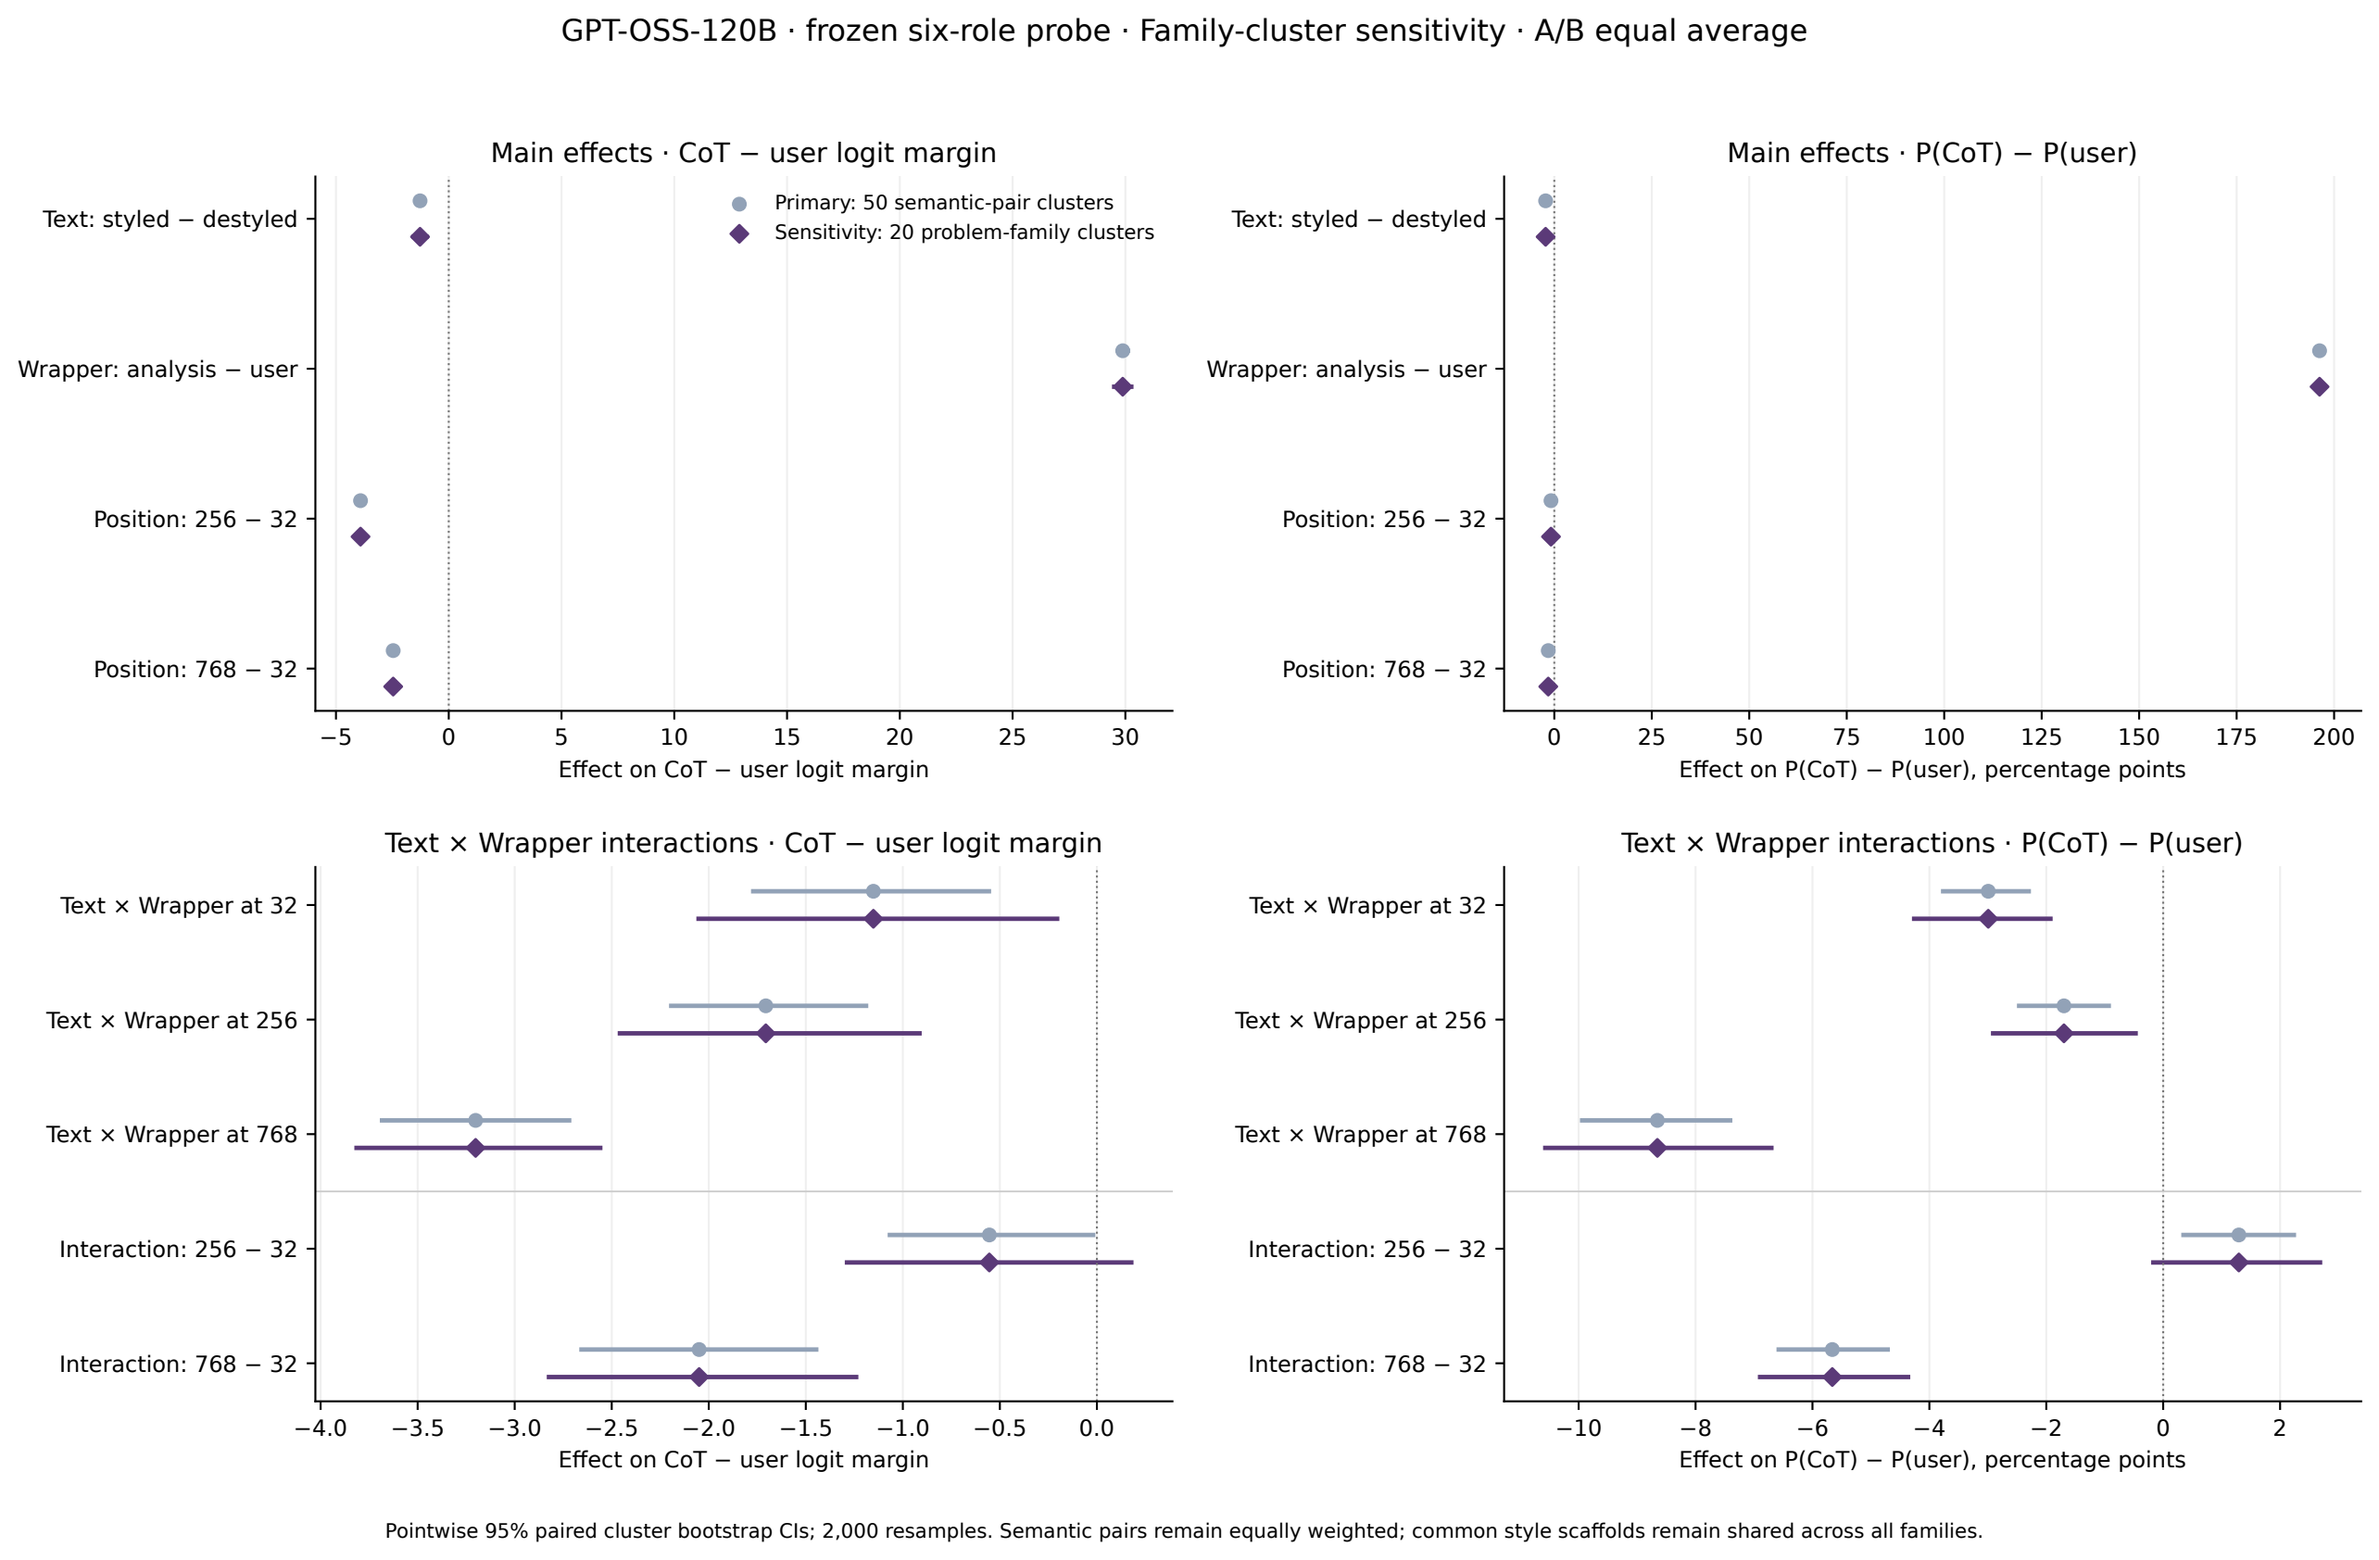
GPT-OSS-120B · frozen six-role probe · Family-cluster sensitivity · A/B equal average
Main effects · CoT − user logit margin
Primary: 50 semantic-pair clusters
Sensitivity: 20 problem-family clusters
Text: styled − destyled
Wrapper: analysis − user
Position: 256 − 32
Position: 768 − 32
−5
0
5
10
15
20
25
30
Effect on CoT − user logit margin
Main effects · P(CoT) − P(user)
Text: styled − destyled
Wrapper: analysis − user
Position: 256 − 32
Position: 768 − 32
0
25
50
75
100
125
150
175
200
Effect on P(CoT) − P(user), percentage points
Text × Wrapper interactions · CoT − user logit margin
Text × Wrapper at 32
Text × Wrapper at 256
Text × Wrapper at 768
Interaction: 256 − 32
Interaction: 768 − 32
−4.0
−3.5
−3.0
−2.5
−2.0
−1.5
−1.0
−0.5
0.0
Effect on CoT − user logit margin
Text × Wrapper interactions · P(CoT) − P(user)
Text × Wrapper at 32
Text × Wrapper at 256
Text × Wrapper at 768
Interaction: 256 − 32
Interaction: 768 − 32
−10
−8
−6
−4
−2
0
2
Effect on P(CoT) − P(user), percentage points
Pointwise 95% paired cluster bootstrap CIs; 2,000 resamples. Semantic pairs remain equally weighted; common style scaffolds remain shared across all families.