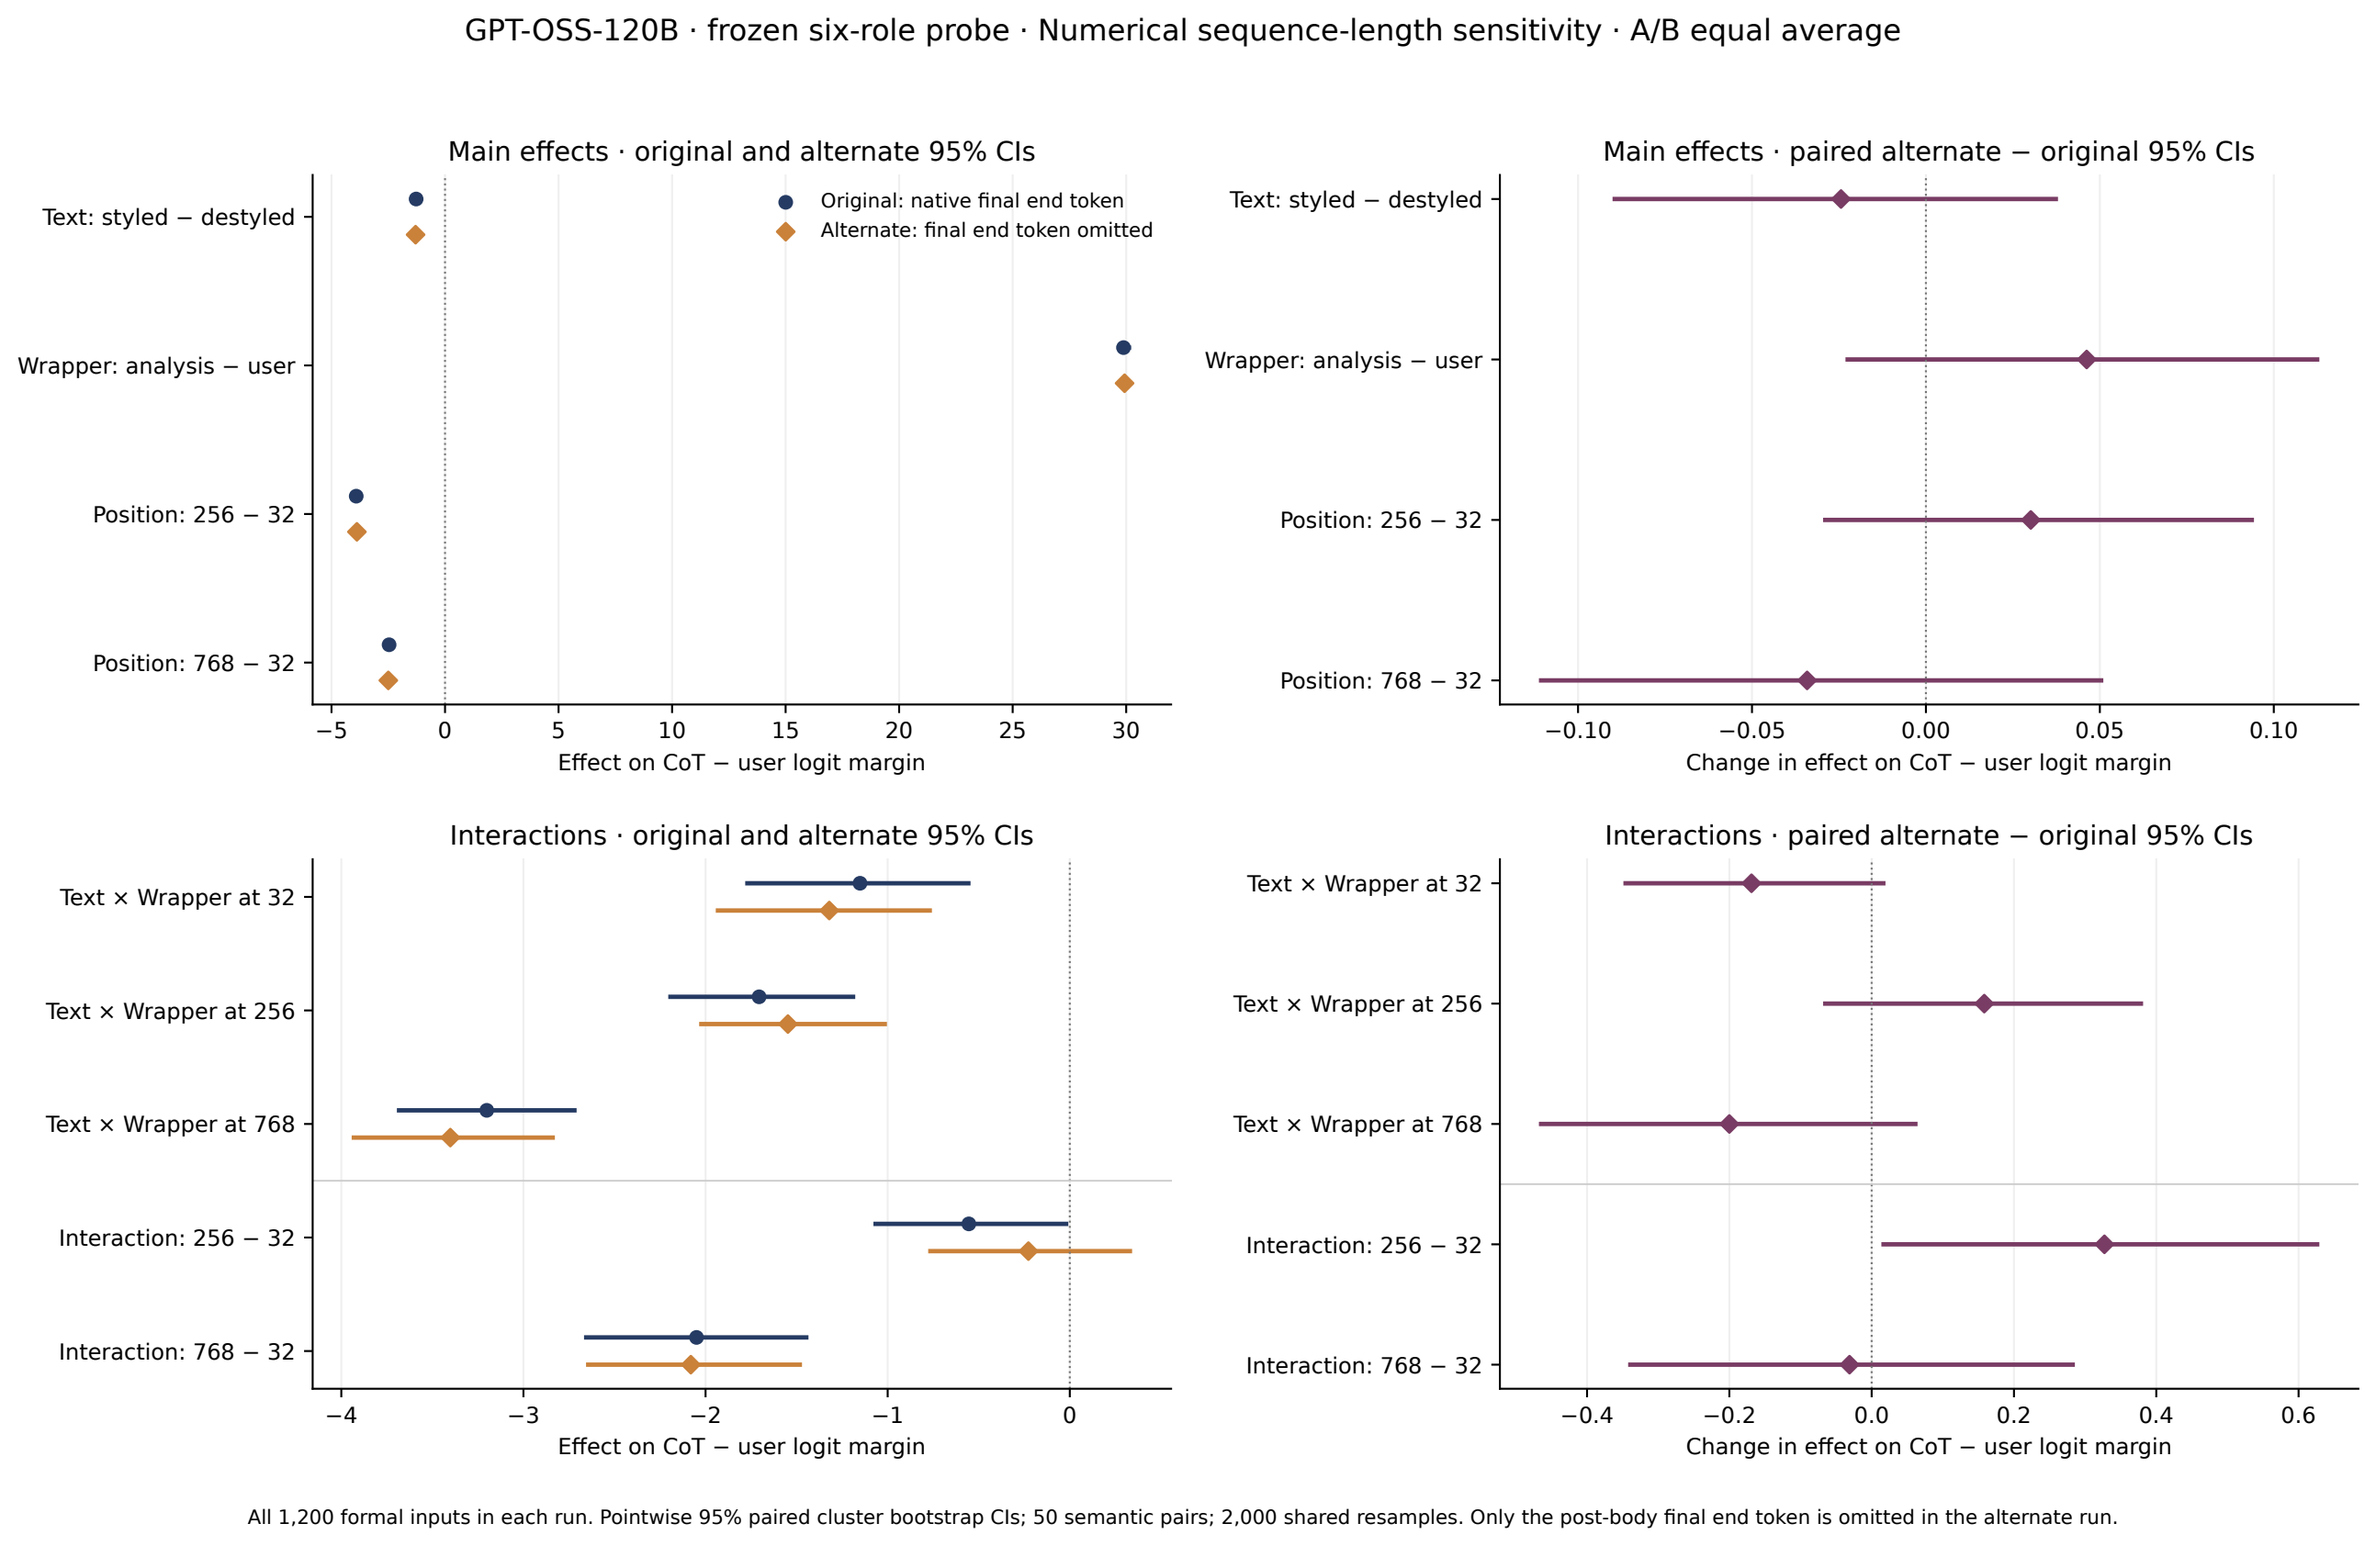
GPT-OSS-120B · frozen six-role probe · Numerical sequence-length sensitivity · A/B equal average
Main effects · original and alternate 95% CIs
Original: native final end token
Alternate: final end token omitted
Text: styled − destyled
Wrapper: analysis − user
Position: 256 − 32
Position: 768 − 32
−5
0
5
10
15
20
25
30
Effect on CoT − user logit margin
Main effects · paired alternate − original 95% CIs
Text: styled − destyled
Wrapper: analysis − user
Position: 256 − 32
Position: 768 − 32
−0.10
−0.05
0.00
0.05
0.10
Change in effect on CoT − user logit margin
Interactions · original and alternate 95% CIs
Text × Wrapper at 32
Text × Wrapper at 256
Text × Wrapper at 768
Interaction: 256 − 32
Interaction: 768 − 32
−4
−3
−2
−1
0
Effect on CoT − user logit margin
Interactions · paired alternate − original 95% CIs
Text × Wrapper at 32
Text × Wrapper at 256
Text × Wrapper at 768
Interaction: 256 − 32
Interaction: 768 − 32
−0.4
−0.2
0.0
0.2
0.4
0.6
Change in effect on CoT − user logit margin
All 1,200 formal inputs in each run. Pointwise 95% paired cluster bootstrap CIs; 50 semantic pairs; 2,000 shared resamples. Only the post-body final end token is omitted in the alternate run.