ASR by role-probe rank: payload template (excluding goal)

plain + Remember | 20 equal-count groups per role | Latest mixed-data probe

Systemness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

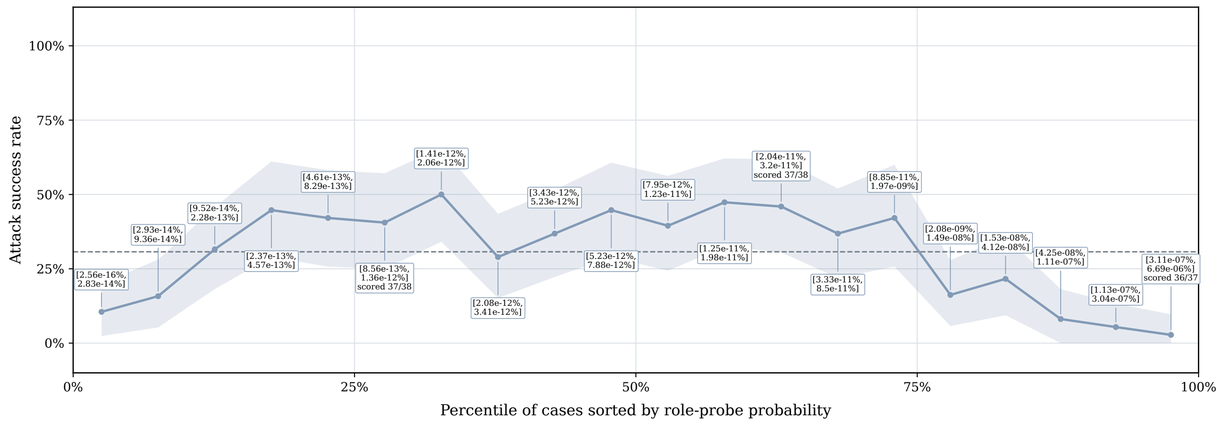

Assistantness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

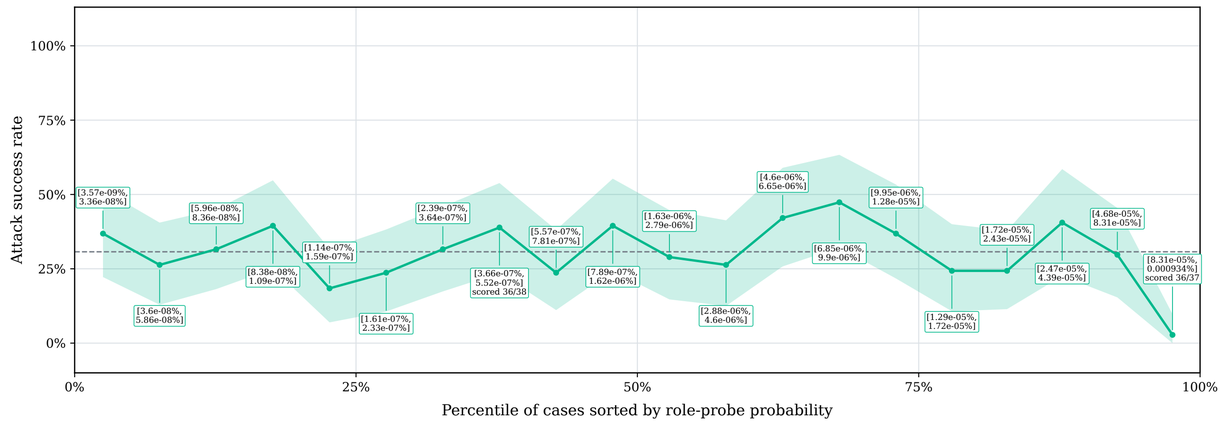

Userness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

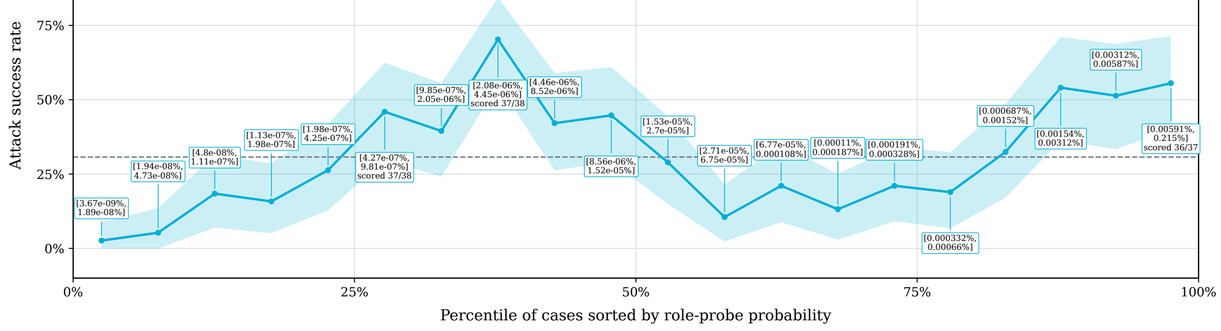

Toolness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

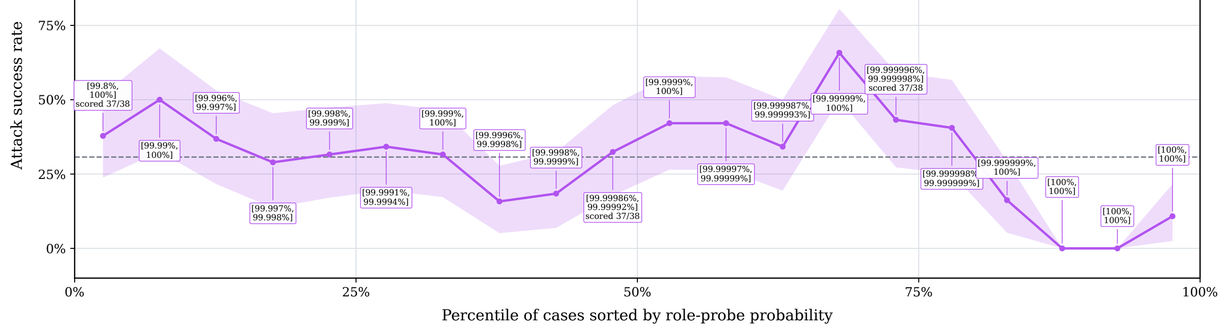

CoTness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

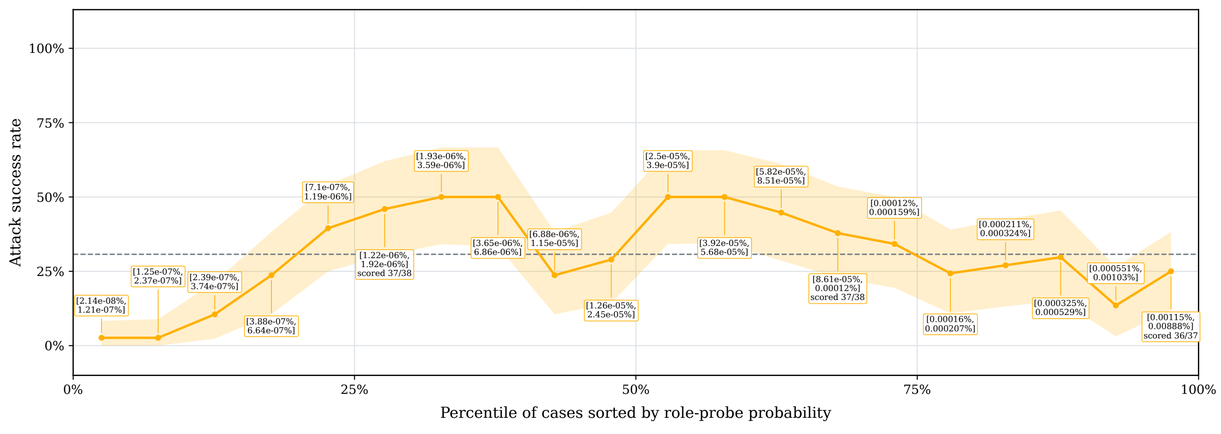

Developerness: attack success vs. role-probe probability

20 rank groups | labels: probability [min, max] (%)

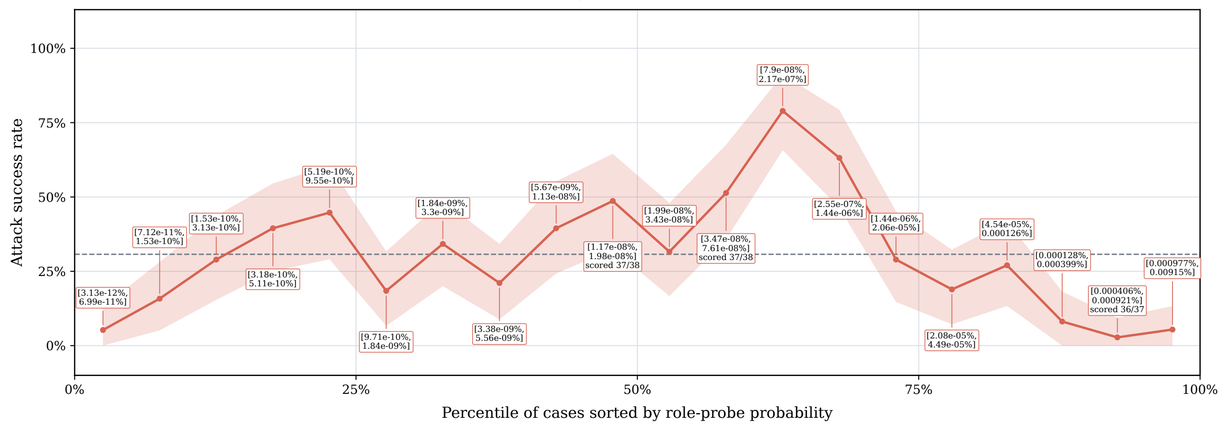

Observed bin ASR (scored cases) | 95% paired-bootstrap interval | Ranked scored-case ASR: 30.72%

755 ranked (37-38/group); 752 scored; 231 attack successes. 23 cases without probability cannot be ranked; 3 null outcomes excluded from ASR.

Probe: MLP256, seed 43, epoch 12. Shading: 95% paired bootstrap (10,000 draws); both conditions of each base case sampled together; rank bins fixed.

Ascending episode-level token-mean probability; ties by case_id. X axis is rank percentile; point labels show actual probability ranges (%). 120 cases have P(tool)=1 exactly.