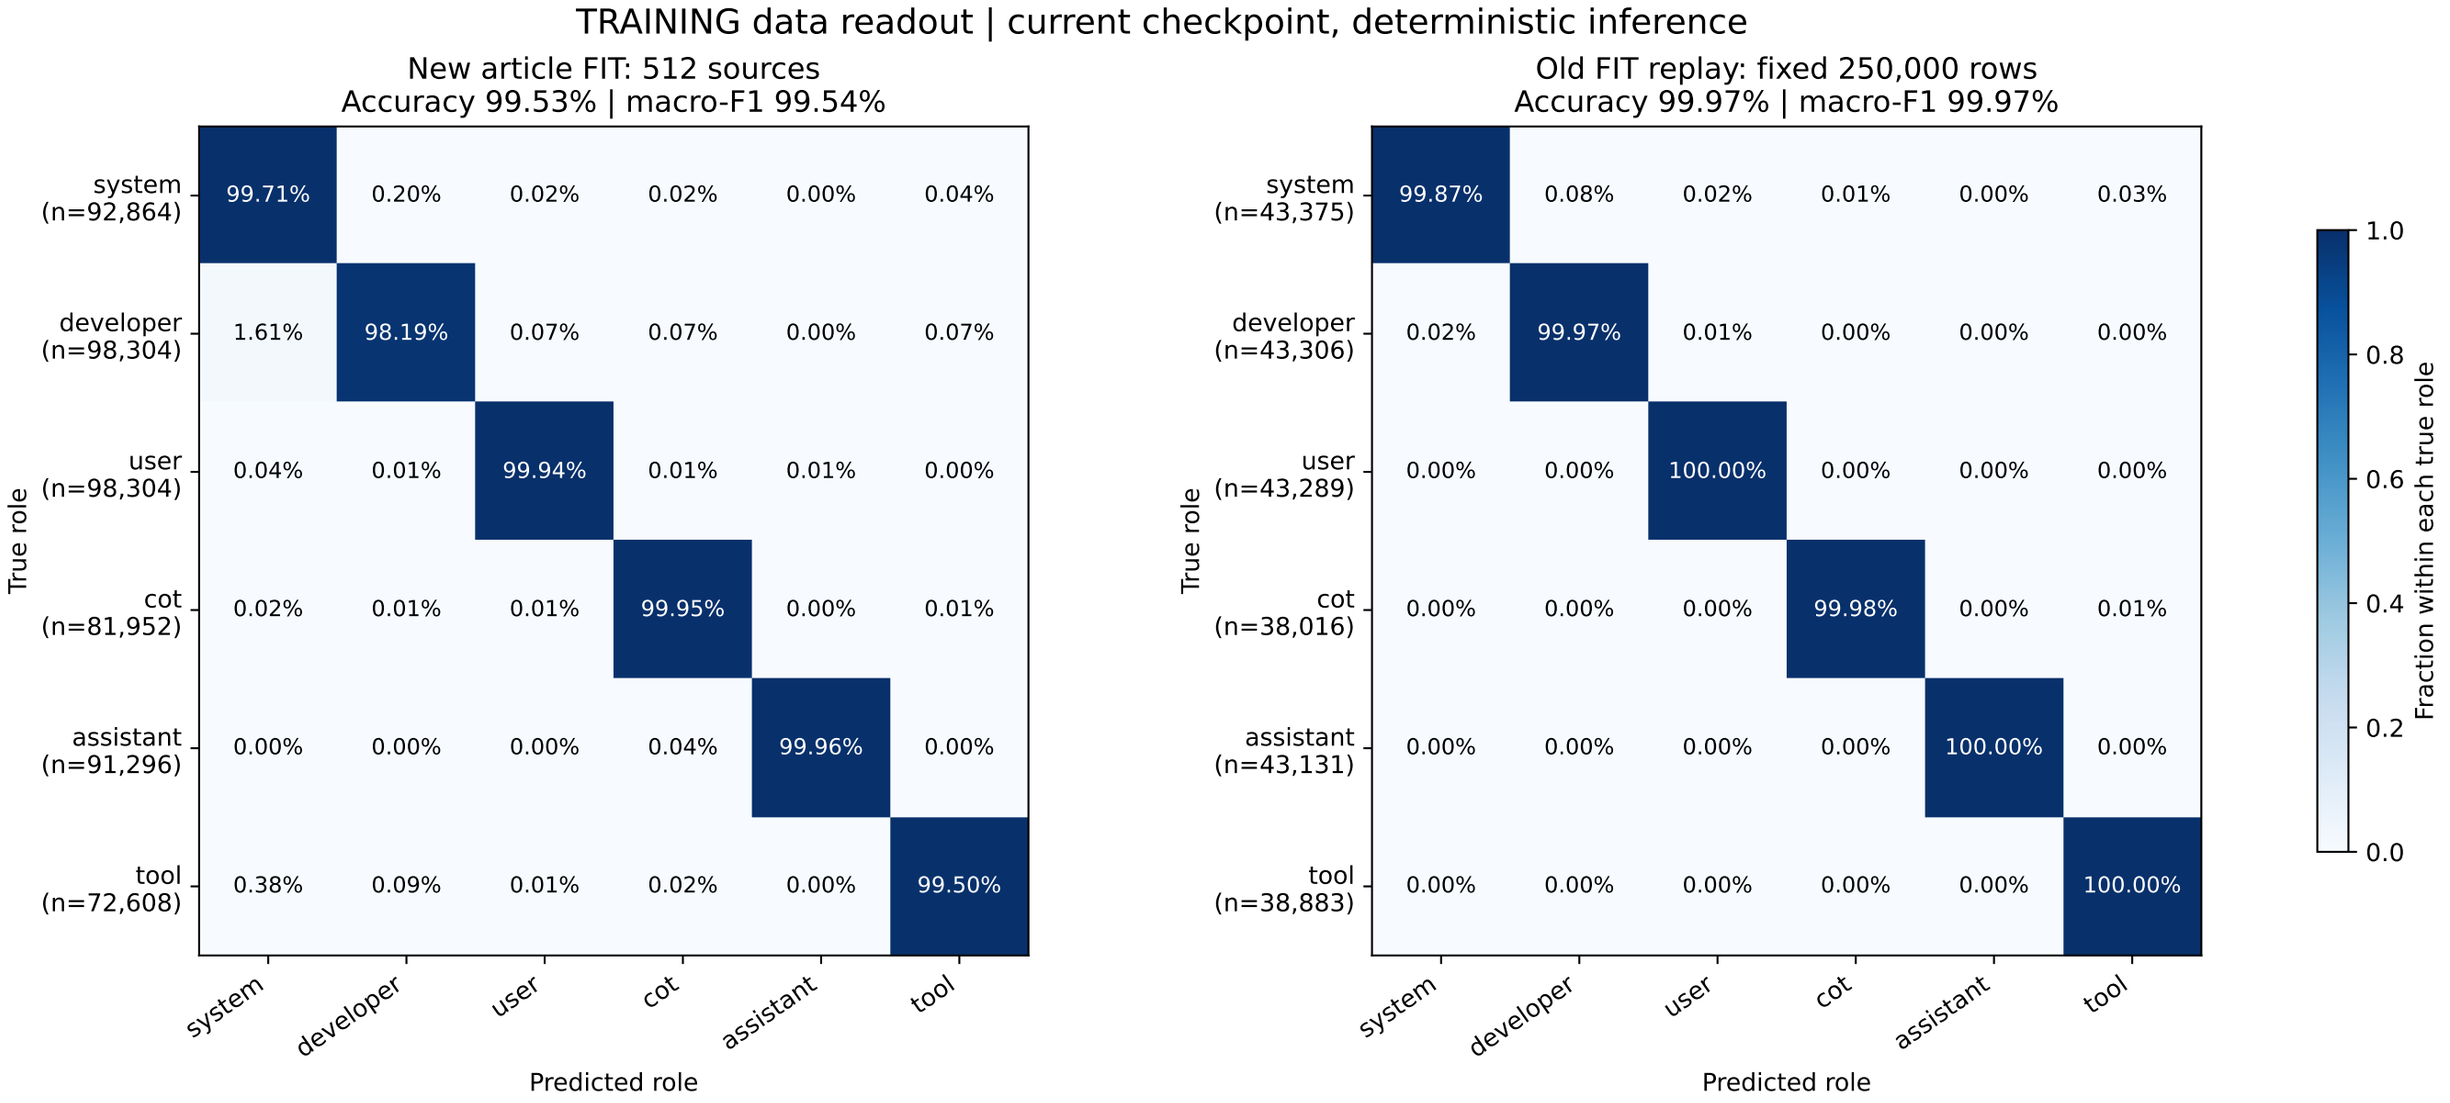
TRAINING data readout | current checkpoint, deterministic inference
New article FIT: 512 sources
Accuracy 99.53% | macro-F1 99.54%
Old FIT replay: fixed 250,000 rows
Accuracy 99.97% | macro-F1 99.97%
True role
system (n=92,864)
developer (n=98,304)
user (n=98,304)
cot (n=81,952)
assistant (n=91,296)
tool (n=72,608)
99.71% 0.20% 0.02% 0.02% 0.00% 0.04%
1.61% 98.19% 0.07% 0.07% 0.00% 0.07%
0.04% 0.01% 99.94% 0.01% 0.01% 0.00%
0.02% 0.01% 0.01% 99.95% 0.00% 0.01%
0.00% 0.00% 0.00% 0.04% 99.96% 0.00%
0.38% 0.09% 0.01% 0.02% 0.00% 99.50%
system
developer
user
cot
assistant
tool
Predicted role
True role
system (n=43,375)
developer (n=43,306)
user (n=43,289)
cot (n=38,016)
assistant (n=43,131)
tool (n=38,883)
99.87% 0.08% 0.02% 0.01% 0.00% 0.03%
0.02% 99.97% 0.01% 0.00% 0.00% 0.00%
0.00% 0.00% 100.00% 0.00% 0.00% 0.00%
0.00% 0.00% 0.00% 99.98% 0.00% 0.01%
0.00% 0.00% 0.00% 0.00% 100.00% 0.00%
0.00% 0.00% 0.00% 0.00% 0.00% 100.00%
system
developer
user
cot
assistant
tool
Predicted role
Fraction within each true role
1.0
0.8
0.6
0.4
0.2
0.0
Article-derived synthetic role contexts; article-body positions only. No AgentDojo/gardening/real-dialogue dataset.
Row-normalized; rounded percentages. Current scratch epoch11 is an INELIGIBLE research candidate.